Exam Author

Premium exam material

Get certification quickly with the Exam Author Premium exam material.
Everything you need to prepare, learn & pass your certification exam easily. Lifetime free updates
First attempt guaranteed success.
https://www.ExamAuthor.com

Microsoft

(DP-700)

Implementing Data Engineering Solutions Using Microsoft Fabric
(beta)

Total: **141 Questions**

Link: https://examauthor.com/papers/microsoft/dp-700

Question: 1 Exam Author

Case Study -
This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.
To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.
At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.
To start the case study -
To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions.
Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.
Overview. Company Overview -
Contoso, Ltd. is an online retail company that wants to modernize its analytics platform by moving to Fabric. The company plans to begin using Fabric for marketing analytics.
Overview. IT Structure -
The company's IT department has a team of data analysts and a team of data engineers that use analytics systems.
The data engineers perform the ingestion, transformation, and loading of data. They prefer to use Python or SQL to transform the data.
The data analysts query data and create semantic models and reports. They are qualified to write queries in Power Query and T-SQL.
Existing Environment. Fabric -
Contoso has an F64 capacity named Cap1. All Fabric users are allowed to create items.
Contoso has two workspaces named WorkspaceA and WorkspaceB that currently use Pro license mode.
Existing Environment. Source Systems
Contoso has a point of sale (POS) system named POS1 that uses an instance of SQL Server on Azure Virtual Machines in the same Microsoft Entra tenant as Fabric. The host virtual machine is on a private virtual network that has public access blocked. POS1 contains all the sales transactions that were processed on the company's website.
The company has a software as a service (SaaS) online marketing app named MAR1. MAR1 has seven entities. The entities contain data that relates to email open rates and interaction rates, as well as website interactions. The data can be exported from MAR1 by calling REST APIs. Each entity has a different endpoint.
Contoso has been using MAR1 for one year. Data from prior years is stored in Parquet files in an Amazon Simple Storage Service (Amazon S3) bucket. There are 12 files that range in size from 300 MB to 900 MB and relate to email interactions.
Existing Environment. Product Data
POS1 contains a product list and related data. The data comes from the following three tables:
Products -
ProductCategories -
ProductSubcategories -
In the data, products are related to product subcategories, and subcategories are related to product categories.
Existing Environment. Azure -
Contoso has a Microsoft Entra tenant that has the following mail-enabled security groups:
DataAnalysts: Contains the data analysts
DataEngineers: Contains the data engineers
Contoso has an Azure subscription.
The company has an existing Azure DevOps organization and creates a new project for repositories that relate to Fabric.
Existing Environment. User Problems
The VP of marketing at Contoso requires analysis on the effectiveness of different types of email content. It typically takes a week to manually compile and analyze the data. Contoso wants to reduce the time to less than one day by using Fabric.
The data engineering team has successfully exported data from MAR1. The team experiences transient connectivity errors, which causes the data exports to fail.
Requirements. Planned Changes -
Contoso plans to create the following two lakehouses:
Lakehouse1: Will store both raw and cleansed data from the sources
Lakehouse2: Will serve data in a dimensional model to users for analytical queries

Additional items will be added to facilitate data ingestion and transformation.
Contoso plans to use Azure Repos for source control in Fabric.
Requirements. Technical Requirements
The new lakehouses must follow a medallion architecture by using the following three layers: bronze, silver, and gold. There will be extensive data cleansing required to populate the MAR1 data in the silver layer, including deduplication, the handling of missing values, and the standardizing of capitalization.
Each layer must be fully populated before moving on to the next layer. If any step in populating the lakehouses fails, an email must be sent to the data engineers.
Data imports must run simultaneously, when possible.
The use of email data from the Amazon S3 bucket must meet the following requirements:
Minimize egress costs associated with cross-cloud data access.
Prevent saving a copy of the raw data in the lakehouses.
Items that relate to data ingestion must meet the following requirements:
The items must be source controlled alongside other workspace items.
Ingested data must land in the bronze layer of Lakehouse1 in the Delta format.
No changes other than changes to the file formats must be implemented before the data lands in the bronze layer.
Development effort must be minimized and a built-in connection must be used to import the source data.
In the event of a connectivity error, the ingestion processes must attempt the connection again.
Lakehouses, data pipelines, and notebooks must be stored in WorkspaceA. Semantic models, reports, and dataflows must be stored in WorkspaceB.
Once a week, old files that are no longer referenced by a Delta table log must be removed.
Requirements. Data Transformation
In the POS1 product data, ProductID values are unique. The product dimension in the gold layer must include only active products from product list. Active products are identified by an IsActive value of 1.
Some product categories and subcategories are NOT assigned to any product. They are NOT analytically relevant and must be omitted from the product dimension in the gold layer.
Requirements. Data Security -
Security in Fabric must meet the following requirements:
The data engineers must have read and write access to all the lakehouses, including the underlying files.
The data analysts must only have read access to the Delta tables in the gold layer.
The data analysts must NOT have access to the data in the bronze and silver layers.
The data engineers must be able to commit changes to source control in WorkspaceA.
You need to ensure that the data analysts can access the gold layer lakehouse.
What should you do?

A.Add the DataAnalyst group to the Viewer role for WorkspaceA.

B.Share the lakehouse with the DataAnalysts group and grant the Build reports on the default semantic model permission.

C.Share the lakehouse with the DataAnalysts group and grant the Read all SQL Endpoint data permission.

D.Share the lakehouse with the DataAnalysts group and grant the Read all Apache Spark permission.

Answer: C

Explanation:

C: Share the lakehouse with the DataAnalysts group and grant the Read all data permission.This approach ensures that data analysts have the necessary read access to the Delta tables in the gold layer, aligning with the requirement that they should not have access to data in the bronze and silver layers.

By granting Read all SQL Endpoint data permission, the analysts get the necessary and sufficient access to query the gold layer data while adhering to the principle of least privilege.

Question: 2

You have a Fabric workspace.
You have semi-structured data.
You need to read the data by using T-SQL, KQL, and Apache Spark. The data will only be written by using Spark.
What should you store the data?

A.a lakehouse
B.an eventhouse
C.a datamart
D.a warehouse

Answer: B

Explanation:

B. an eventhouse .

- T-SQL, KQL and Spark can be used on Eventhouse

- T-SQL & Spark can be used on Lakehouse

An eventhouse suggests the focus is on storing event-based or streaming data. While the data itself is semi-structured and written using Apache Spark, the choice of an eventhouse aligns with scenarios where real-time ingestion and analysis of data streams are required. Additionally, an eventhouse is optimized for applications handling high-frequency data events, making it suitable for Spark-based write operations and enabling the integration of T-SQL, KQL, and Spark query capabilities.

Question: 3

You have a Fabric workspace that contains a warehouse named Warehouse1.
You have an on-premises Microsoft SQL Server database named Database1 that is accessed by using an on-premises data gateway.
You need to copy data from Database1 to Warehouse1.
Which item should you use?

A.a Dataflow Gen1 dataflow
B.a data pipeline
C.a KQL queryset
D.a notebook

Answer: B

Explanation:

B: **a data pipeline.**

The pipeline in Fabric works same way as pipeline in Data Factory. It supports the use of self-hosted Integrated Runtime (IR) which is used to establish connection between a private network (on-prem) and Azure cloud.

A data pipeline is the most suitable tool for moving data between different sources and destinations. In this case, you need to copy data from your on-premises Microsoft SQL Server database (Database1) to your Fabric warehouse (Warehouse1). A data pipeline can efficiently handle this task by allowing you to define and manage the data transfer process.

Question: 4

Exam Author

You have a Fabric workspace that contains a warehouse named Warehouse1.
You have an on-premises Microsoft SQL Server database named Database1 that is accessed by using an on-premises data gateway.
You need to copy data from Database1 to Warehouse1.
Which item should you use?

A.an Apache Spark job definition

B.a data pipeline

C.a Dataflow Gen1 dataflow

D.an eventstream

Answer: B

Explanation:

B: **a data pipeline.**

A data pipeline is specifically designed for orchestrating and automating data movement tasks between different sources and destinations. Here's why a data pipeline is the best choice for copying data from your on-premises Microsoft SQL Server database (Database1) to your Fabric warehouse (Warehouse1)

Data pipelines in Microsoft Fabric are designed to facilitate the movement and transformation of data between various sources and destinations. In this scenario, a data pipeline can be configured to copy data from the on-premises SQL Server database to the Fabric warehouse, utilizing the on-premises data gateway for secure connectivity.

Question: 5

Exam Author

You have a Fabric F32 capacity that contains a workspace. The workspace contains a warehouse named DW1 that is modelled by using MD5 hash surrogate keys.
DW1 contains a single fact table that has grown from 200 million rows to 500 million rows during the past year.
You have Microsoft Power BI reports that are based on Direct Lake. The reports show year-over-year values.
Users report that the performance of some of the reports has degraded over time and some visuals show errors.
You need to resolve the performance issues. The solution must meet the following requirements:
Provide the best query performance.
Minimize operational costs.
Which should you do?

A.Change the MD5 hash to SHA256.

B.Increase the capacity.

C.Enable V-Order.

D.Modify the surrogate keys to use a different data type.

E.Create views.

Answer: C

Explanation:

C. Enable V-Order

Why Enable V-Order is correct

Your scenario highlights:

Large fact table growth (200M → 500M rows)

Direct Lake Power BI reports

Performance degradation and visual errors

Requirement for:

✔ Best query performance

✔ Minimal operational cost

V-Order:

Optimizes the physical layout of Delta tables.

Improves columnar compression and scan efficiency.

Significantly boosts Direct Lake query performance.

Requires no additional capacity cost.

So it directly addresses performance while keeping costs low.

Question: 6 **Exam Author**

HOTSPOT -
You have a Fabric workspace that contains a warehouse named DW1. DW1 contains the following tables and columns.

Table name	Column name	Description
SalesOrderDetail	ProductID	Contains the product ID of the ordered product
SalesOrderDetail	ModifiedDate	Contains the date of an order
SalesOrderDetail	OrderQty	Contains the order quantity
Product	ProductID	Contains the unique ID of a product
Product	Name	Contains a product name

You need to create an output that presents the summarized values of all the order quantities by year and product.
The results must include a summary of the order quantities at the year level for all the products.
How should you complete the code? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

[▼] (SO.ModifiedDate) AS OrderDate

- SELECT CAST
- SELECT CONVERT
- SELECT YEAR

```
    ,P.Name AS ProductName
    ,SUM(SO.OrderQty) AS OrderQty
  FROM [dbo].[SalesOrderDetail] SO
  INNER JOIN [dbo].[Product] P
    ON P.ProductID = SO.ProductID
  GROUP BY
```

[▼]

- CUBE(YEAR(SO.ModifiedDate), P.Name)
- GROUPING SETS ((YEAR(SO.ModifiedDate), P.Name), (YEAR(SO.ModifiedDate)))
- ROLLUP(YEAR(SO.ModifiedDate), P.Name)
- YEAR(SO.ModifiedDate), P.Name

```
  ORDER BY OrderDate
```

Answer:

Box1 -> SELECT YEAR || Box2 -> ROLLUP(YEAR(SO.ModifiedDATE), P.Name)

Explanation:

```
[dropdown] (SO.ModifiedDate) AS OrderDate
```

- SELECT CAST
- SELECT CONVERT
- SELECT YEAR

```
    ,P.Name AS ProductName
    ,SUM(SO.OrderQty) AS OrderQty
  FROM [dbo].[SalesOrderDetail] SO
  INNER JOIN [dbo].[Product] P
    ON P.ProductID = SO.ProductID
  GROUP BY
```

- CUBE(YEAR(SO.ModifiedDate), P.Name)
- GROUPING SETS ((YEAR(SO.ModifiedDate), P.Name), (YEAR(SO.ModifiedDate)))
- ROLLUP(YEAR(SO.ModifiedDate), P.Name)
- YEAR(SO.ModifiedDate), P.Name

```
  ORDER BY OrderDate
```

To address the requirements for summarizing order quantities by year and product, with a summary at the year level for all products, let's analyze the SQL solution:

Explanation of Choices:

1. `SELECT YEAR`: Extracts the year from the `ModifiedDate` column, which is essential for grouping data by year.
2. `ROLLUP(YEAR(SO.ModifiedDate), P.Name)`:
   - Adds subtotals at each hierarchical level of grouping (e.g., yearly totals across all products).
   - Ensures the result includes summaries for both each product by year and all products within a year.

Completed Query:

sql Copy Edit

```sql
SELECT YEAR(SO.ModifiedDate) AS OrderDate,
       P.Name AS ProductName,
       SUM(SO.OrderQty) AS OrderQty
FROM [dbo].[SalesOrderDetail] SO
INNER JOIN [dbo].[Product] P
    ON P.ProductID = SO.ProductID
GROUP BY ROLLUP(YEAR(SO.ModifiedDate), P.Name)
ORDER BY OrderDate;
```

Key Details:

The use of ROLLUP ensures compliance with the requirement for summarized values at different grouping levels.
SUM(SO.OrderQty) calculates the total order quantities.

Question: 7

You have a Fabric workspace that contains a lakehouse named Lakehouse1. Data is ingested into Lakehouse1 as one flat table. The table contains the following columns.

Name	Description
TransactionID	Contains a unique ID for each transaction
Date	Contains the date of a transaction
ProductID	Contains a unique ID for each product
ProductColor	Contains a descriptive attribute that describes the color of each product
ProductName	Contains a unique name for each product
SalesAmount	Contains the sales amount of a transaction

You plan to load the data into a dimensional model and implement a star schema. From the original flat table, you create two tables named FactSales and DimProduct. You will track changes in DimProduct.
You need to prepare the data.
Which three columns should you include in the DimProduct table? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A.Date
B.ProductName
C.ProductColor
D.TransactionID
E.SalesAmount
F.ProductID

Answer: BCF

Explanation:

B. **ProductName:** This attribute describes the product and is crucial for understanding and analyzing the data related to each product.

C. **ProductColor:** This attribute provides additional information about the product, which can be useful for analysis, reporting, and segmentation.

F. **ProductID:** This is the unique identifier for each product and serves as the primary key for the DimProduct table. It's essential for establishing the relationship between the FactSales table and the DimProduct table.

Question: 8

You have a Fabric workspace named Workspace1 that contains a notebook named Notebook1.
In Workspace1, you create a new notebook named Notebook2.
You need to ensure that you can attach Notebook2 to the same Apache Spark session as Notebook1.
What should you do?

A.Enable high concurrency for notebooks.

B.Enable dynamic allocation for the Spark pool.

C.Change the runtime version.

D.Increase the number of executors.

Answer: A

Explanation:

A.Enable high concurrency for notebooks: High concurrency allows multiple notebooks to share the same Apache Spark session. This setting ensures that different notebooks can run simultaneously within the same session, facilitating collaboration and efficient resource usage.

Question: 9 **Exam Author**

You have a Fabric workspace named Workspace1 that contains a lakehouse named Lakehouse1. Lakehouse1 contains the following tables:
Orders -
Customer -
Employee -
The Employee table contains Personally Identifiable Information (PII).
A data engineer is building a workflow that requires writing data to the Customer table, however, the user does NOT have the elevated permissions required to view the contents of the Employee table.
You need to ensure that the data engineer can write data to the Customer table without reading data from the Employee table.
Which three actions should you perform? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A.Share Lakehouse1 with the data engineer.

B.Assign the data engineer the Contributor role for Workspace2.

C.Assign the data engineer the Viewer role for Workspace2.

D.Assign the data engineer the Contributor role for Workspace1.

E.Migrate the Employee table from Lakehouse1 to Lakehouse2.

F.Create a new workspace named Workspace2 that contains a new lakehouse named Lakehouse2.

G.Assign the data engineer the Viewer role for Workspace1.

Answer: DEF

Explanation:

D. **Assign the data engineer the Contributor role for Workspace1**:

Assigning the Contributor role to the data engineer for Workspace1 grants them the necessary permissions to write data to the Customer table in Lakehouse1. However, since the data engineer does not have elevated permissions to view the Employee table, they won't be able to access its content.

E. **Migrate the Employee table from Lakehouse1 to Lakehouse2:**

Moving the Employee table, which contains Personally Identifiable Information (PII), to a separate Lakehouse2 helps ensure that the data engineer cannot accidentally or intentionally access it. This action keeps sensitive data segregated from the data engineer's operational environment.

F. **Create a new workspace named Workspace2 that contains a new lakehouse named Lakehouse2:**

By creating a new workspace and lakehouse for the Employee table, you further isolate the sensitive data. The data engineer can still perform their tasks in Workspace1 without accessing Workspace2, ensuring secure data handling and compliance with privacy requirements.

Question: 10 **Exam Author**

You have a Fabric warehouse named DW1. DW1 contains a table that stores sales data and is used by multiple sales representatives.
You plan to implement row-level security (RLS).
You need to ensure that the sales representatives can see only their respective data.
Which warehouse object do you require to implement RLS?

A.STORED PROCEDURE
B.CONSTRAINT
C.SCHEMA
D.FUNCTION

Answer: D

Explanation:

To implement Row-Level Security (RLS) in a Fabric warehouse like DW1, need to use a FUNCTION to define the filtering logic. Specifically, a user-defined function (UDF) is created and associated with the RLS policy to determine which rows each user can access.

Reference:

https://learn.microsoft.com/en-us/fabric/data-warehouse/tutorial-row-level-security#2-define-security-policies

Question: 11 **Exam Author**

HOTSPOT -
You have a Fabric workspace named Workspace1_DEV that contains the following items:
10 reports
Four notebooks -
Three lakehouses -
Two data pipelines -
Two Dataflow Gen1 dataflows -
Three Dataflow Gen2 dataflows -
Five semantic models that each has a scheduled refresh policy
You create a deployment pipeline named Pipeline1 to move items from Workspace1_DEV to a new workspace named Workspace1_TEST.
You deploy all the items from Workspace1_DEV to Workspace1_TEST.
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Data from the semantic models will be deployed to the target stage.		
The Dataflow Gen1 dataflows will be deployed to the target stage.		
The scheduled refresh policies will be deployed to the target stage.		

Answer: No/Yes/No

Explanation:

Statements	Yes	No
Data from the semantic models will be deployed to the target stage.		
The Dataflow Gen1 dataflows will be deployed to the target stage.		
The scheduled refresh policies will be deployed to the target stage.		

1. Data from the semantic models will be deployed to the target stage.

Answer: No
Semantic models are only deployed to the target stage in the form of **metadata**. The deployment process does not copy actual data; instead, only the structural and configuration metadata (e.g., model schema and measures) is deployed. The target stage will require a refresh to fetch the data into the semantic models.
Reference: Microsoft Learn - Item Properties Copied During Deployment

2.The Dataflow Gen1 dataflows will be deployed to the target stage.

Answer: Yes
Dataflow Gen1 objects are included in the deployment pipeline and are fully deployed to the target stage, including their configurations. This ensures that Dataflow Gen1 pipelines can run in the target environment. The deployment process supports this functionality without requiring a manual configuration.

3.The scheduled refresh policies will be deployed to the target stage.

Answer: No
The deployment process does not copy or deploy **refresh schedules** for datasets, semantic models, or other items. Although metadata for the items is deployed, refresh schedules must be manually recreated or configured in the target stage. This limitation is highlighted in Microsoft's documentation.
Reference: Microsoft Learn - Item Properties Copied During Deployment

Question: 12

Exam Author

You have a Fabric deployment pipeline that uses three workspaces named Dev, Test, and Prod.
You need to deploy an eventhouse as part of the deployment process.
What should you use to add the eventhouse to the deployment process?

A.GitHub Actions
B.a deployment pipeline
C.an Azure DevOps pipeline

Answer: B

Explanation:

B. a deployment pipeline

Microsoft Fabric deployment pipelines are designed for managing the deployment of resources (including eventhouses) across different environments (e.g., Dev, Test, Prod). You use a deployment pipeline to automate and control the deployment process across workspaces, ensuring that the eventhouse (along with other resources) is deployed correctly through the stages of your workflow.

GitHub Actions and Azure DevOps pipelines could also be used for deployment, but the Fabric deployment pipeline is specifically built for this use case within Microsoft Fabric and integrates directly with the workspace and stage management.

Get started with deployment pipelines

https://learn.microsoft.com/en-us/fabric/cicd/deployment-pipelines/get-started-with-deployment-pipelines?tabs=from-fabric%2Cnew-ui

Question: 13

Exam Author

You have a Fabric workspace named Workspace1 that contains a warehouse named Warehouse1.
You plan to deploy Warehouse1 to a new workspace named Workspace2.
As part of the deployment process, you need to verify whether Warehouse1 contains invalid references. The solution must minimize development effort.
What should you use?

A.a database project
B.a deployment pipeline
C.a Python script
D.a T-SQL script

Answer: B

Explanation:

Microsoft Fabric's deployment pipelines provide a built-in mechanism to manage and validate the deployment of artifacts like warehouses. When you use a deployment pipeline to move Warehouse1 from one workspace (Workspace1) to another (Workspace2), the pipeline automatically checks for issues such as invalid references or missing dependencies during the deployment process.

Question: 14 **Exam Author**

You have a Fabric workspace that contains a Real-Time Intelligence solution and an eventhouse.
Users report that from OneLake file explorer, they cannot see the data from the eventhouse.
You enable OneLake availability for the eventhouse.
What will be copied to OneLake?

A.only data added to new databases that are added to the eventhouse

B.only the existing data in the eventhouse

C.no data

D.both new data and existing data in the eventhouse

E.only new data added to the eventhouse

Answer: D

Explanation:

D. both new data and existing data in the eventhouseWhen you enable OneLake availability for the eventhouse, both the existing data and any new data added to the eventhouse will be copied to OneLake. This allows users to access the eventhouse data from OneLake file explorer.

Question: 15 **Exam Author**

You have a Fabric workspace named Workspace1.
You plan to integrate Workspace1 with Azure DevOps.
You will use a Fabric deployment pipeline named deployPipeline1 to deploy items from Workspace1 to higher environment workspaces as part of a medallion architecture. You will run deployPipeline1 by using an API call from an Azure DevOps pipeline.
You need to configure API authentication between Azure DevOps and Fabric.
Which type of authentication should you use?

A.service principal

B.Microsoft Entra username and password

C.managed private endpoint

D.workspace identity

Answer: A

Explanation:

A. service principal.

Service Principal: A service principal is a security identity used by applications, services, and automation tools to access specific Azure resources. It provides a secure way to authenticate and authorize API calls between Azure DevOps and Fabric. By using a service principal, you can grant the necessary permissions to deployPipeline1 to interact with the Fabric workspace (Workspace1) and deploy items to higher environments. This approach ensures secure and managed access without relying on individual user credentials.

Question: 16 **Exam Author**

You have a Google Cloud Storage (GCS) container named storage1 that contains the files shown in the following table.

Name	Size
ProductFile.parquet	8 MB
StoreFile.json	500 MB
TripsFile.csv	99 MB

You have a Fabric workspace named Workspace1 that has the cache for shortcuts enabled. Workspace1 contains a lakehouse named Lakehouse1. Lakehouse1 has the shortcuts shown in the following table.

Name	Source	Last accessed
Products	ProductFile	12 hours ago
Stores	StoreFile	4 hours ago
Trips	TripsFile	48 hours ago

You need to read data from all the shortcuts.
Which shortcuts will retrieve data from the cache?

A.Stores only
B.Products only
C.Stores and Products only
D.Products, Stores, and Trips
E.Trips only
F.Products and Trips only

Answer: C

Explanation:

C. Stores and Products only.

Both are master/reference data, so they retrieve from the cache.

Shortcut caching can be used to reduce egress costs associated with cross-cloud data access. As files are read through an external shortcut, the files are stored in a cache for the Fabric workspace. Subsequent read requests are served from cache rather than the remote storage provider. Cached files have a retention period of 24 hours.

https://learn.microsoft.com/en-us/fabric/onelake/onelake-shortcuts#caching

Question: 17 **Exam Author**

You have a Fabric workspace named Workspace1 that contains an Apache Spark job definition named Job1.

You have an Azure SQL database named Source1 that has public internet access disabled.
You need to ensure that Job1 can access the data in Source1.
What should you create?

A.an on-premises data gateway

B.a managed private endpoint

C.an integration runtime

D.a data management gateway

Answer: B

Explanation:

B. a managed private endpoint.

Managed Private Endpoint: This allows secure and private communication between Azure services without exposing data to the public internet. By creating a managed private endpoint, you can establish a direct connection between the Apache Spark job in Workspace1 and the Azure SQL database (Source1) while keeping public internet access disabled. This approach ensures that data transfer happens securely within the Azure network.

To ensure that Job1 can access the data in Source1, you need to create a managed private endpoint. This will allow the Spark job to securely connect to the Azure SQL database without requiring public internet access.

Question: 18

You have an Azure Data Lake Storage Gen2 account named storage1 and an Amazon S3 bucket named storage2.
You have the Delta Parquet files shown in the following table.

Name	Stored in	Size	Description
ProductFile	storage1	50 MB	Contains a list of products and their details
TripsFile	storage2	2 GB	Contains one month's worth of taxi trip data
StoreFile	storage2	25 MB	Contains a list of stores and their addresses

You have a Fabric workspace named Workspace1 that has the cache for shortcuts enabled. Workspace1 contains a lakehouse named Lakehouse1. Lakehouse1 has the following shortcuts:
A shortcut to ProductFile aliased as Products
A shortcut to StoreFile aliased as Stores
A shortcut to TripsFile aliased as Trips
The data from which shortcuts will be retrieved from the cache?

A.Trips and Stores only

B.Products and Store only

C.Stores only

D.Products only

E.Products, Stores, and Trips

Answer: C

Explanation:

Because according to the current Microsoft Fabric shortcut caching rules:

Shortcut caching is only supported for GCS, S3, and S3-compatible storage —

ADLS Gen2 is not included in the caching support list.

→ This excludes ProductFile (even though it's small, it's on ADLS Gen2).

Files over 1 GB are not cached —

→ This excludes TripsFile (2 GB, S3).

StoreFile is on S3 (supported) and is 25 MB (<1 GB) —

→ Eligible for caching.

Result: Only StoreFile meets the criteria, so the correct choice is C. Stores only.

Question: 19 **Exam Author**

HOTSPOT -
You have a Fabric workspace named Workspace1 that contains the items shown in the following table.

Name	Type
Notebook1	Notebook
Notebook2	Notebook
Lakehouse1	Lakehouse
Pipeline1	Data pipeline
Model1	Semantic model

For Model1, the Keep your Direct Lake data up to date option is disabled.
You need to configure the execution of the items to meet the following requirements:
Notebook1 must execute every weekday at 8:00 AM.
Notebook2 must execute when a file is saved to an Azure Blob Storage container.
Model1 must refresh when Notebook1 has executed successfully.
How should you orchestrate each item? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Notebook1:

- Add Notebook1 to an Apache Spark job definition.
- Add Notebook1 to Pipeline1.
- From Real-Time hub, configure the execution of Notebook1.

Notebook2:

- Add Notebook2 to an Apache Spark job definition.
- Add Notebook2 to Pipeline1.
- From Real-Time hub, configure the execution of Notebook2.

Pipeline1:

- Add Pipeline1 to an Apache Spark job definition.
- Configure the execution of Pipeline1 by using a schedule.
- From Real-Time hub, configure the execution of Pipeline1.

Model1:

- Add Model1 to Pipeline1.
- From Real-Time hub, configure Model1 to refresh.
- Set Keep your Direct Lake data up to date to On.

Answer:

Box1 ->2 || Box2 -> 3 || Box3 -> 2 || Box4 -> 1

Explanation:

Notebook1:

- Add Notebook1 to an Apache Spark job definition.
- Add Notebook1 to Pipeline1.
- From Real-Time hub, configure the execution of Notebook1.

Notebook2:

- Add Notebook2 to an Apache Spark job definition.
- Add Notebook2 to Pipeline1.
- From Real-Time hub, configure the execution of Notebook2.

Pipeline1:

- Add Pipeline1 to an Apache Spark job definition.
- Configure the execution of Pipeline1 by using a schedule.
- From Real-Time hub, configure the execution of Pipeline1.

Model1:

- Add Model1 to Pipeline1.
- From Real-Time hub, configure Model1 to refresh.
- Set Keep your Direct Lake data up to date to On.

1. **Notebook1**:
   - **Selection**: Add Notebook1 to Pipeline1.
   - **Justification**: To schedule Notebook1 for execution every weekday at 8:00 AM, adding it to a pipeline that allows scheduled execution is the correct choice.
2. **Notebook2**:
   - **Selection**: From Real-Time hub, configure the execution of Notebook2.
   - **Justification**: The requirement states that Notebook2 should execute when a file is saved to an Azure Blob Storage container. The Real-Time hub provides event-driven execution, making this the correct option.
3. **Pipeline1**:
   - **Selection**: Configure the execution of Pipeline1 by using a schedule.
   - **Justification**: Pipeline1 likely orchestrates multiple tasks, and scheduling its execution ensures that it runs automatically according to a predefined timetable.
4. **Model1**:
   - **Selection**: Add Model1 to Pipeline1.
   - **Justification**: Since Model1 needs to refresh after Notebook1 executes successfully, including it in the same pipeline ensures proper orchestration and dependencies.

Question: 20 Exam Author

Your company has a sales department that uses two Fabric workspaces named Workspace1 and Workspace2.
The company decides to implement a domain strategy to organize the workspaces.
You need to ensure that a user can perform the following tasks:
Create a new domain for the sales department.
Create two subdomains: one for the east region and one for the west region.
Assign Workspace1 to the east region subdomain.
Assign Workspace2 to the west region subdomain.
The solution must follow the principle of least privilege.
Which role should you assign to the user?

A.workspace Admin

B.domain admin

C.domain contributor

D.Fabric admin

Answer: D

Explanation:

Fabric Admin: Possesses the highest level of permissions within the Fabric environment, enabling the creation of domains and subdomains, as well as the assignment of resources to those subdomains.

Question: 21 Exam Author

You have a Fabric workspace named Workspace1 that contains a warehouse named DW1 and a data pipeline named Pipeline1.
You plan to add a user named User3 to Workspace1.
You need to ensure that User3 can perform the following actions:
View all the items in Workspace1.
Update the tables in DW1.
The solution must follow the principle of least privilege.
You already assigned the appropriate object-level permissions to DW1.
Which workspace role should you assign to User3?

A.Admin

B.Member

C.Viewer

D.Contributor

Answer: D

Explanation:

D. Contributor is the correct answer.

Requirements:

✔ View all items in the workspace

✔ Update tables in DW1

✔ Least privilege

✔ Object-level permissions on DW1 already assigned

Question: 22 **Exam Author**

You have a Fabric capacity that contains a workspace named Workspace1. Workspace1 contains a lakehouse named Lakehouse1, a data pipeline, a notebook, and several Microsoft Power BI reports.
A user named User1 wants to use SQL to analyze the data in Lakehouse1.
You need to configure access for User1. The solution must meet the following requirements:
Provide User1 with read access to the table data in Lakehouse1.
Prevent User1 from using Apache Spark to query the underlying files in Lakehouse1.
Prevent User1 from accessing other items in Workspace1.
What should you do?

A.Share Lakehouse1 with User1 directly and select Read all SQL endpoint data.

B.Assign User1 the Viewer role for Workspace1. Share Lakehouse1 with User1 and select Read all SQL endpoint data.

C.Share Lakehouse1 with User1 directly and select Build reports on the default semantic model.

D.Assign User1 the Member role for Workspace1. Share Lakehouse1 with User1 and select Read all SQL endpoint data.

Answer: A

Explanation:

A. Share Lakehouse1 with User1 directly and select Read all SQL endpoint data.

Share Lakehouse1 with User1 directly and select Read all SQL endpoint data: This approach grants User1 read access specifically to the table data in Lakehouse1 through the SQL endpoint, without giving them broader permissions in Workspace1 or access to other items. By directly sharing Lakehouse1 and selecting the "Read all SQL endpoint data" option, you ensure User1 can use SQL to analyze the data while preventing them from using Apache Spark to query the underlying files.

Question: 23 **Exam Author**

DRAG DROP -
You are implementing the following data entities in a Fabric environment:
Entity1: Available in a lakehouse and contains data that will be used as a core organization entity
Entity2: Available in a semantic model and contains data that meets organizational standards
Entity3: Available in a Microsoft Power BI report and contains data that is ready for sharing and reuse
Entity4: Available in a Power BI dashboard and contains approved data for executive-level decision making
Your company requires that specific governance processes be implemented for the data.
You need to apply endorsement badges to the entities based on each entity’s use case.
Which badge should you apply to each entity? To answer, drag the appropriate badges the correct entities. Each badge may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

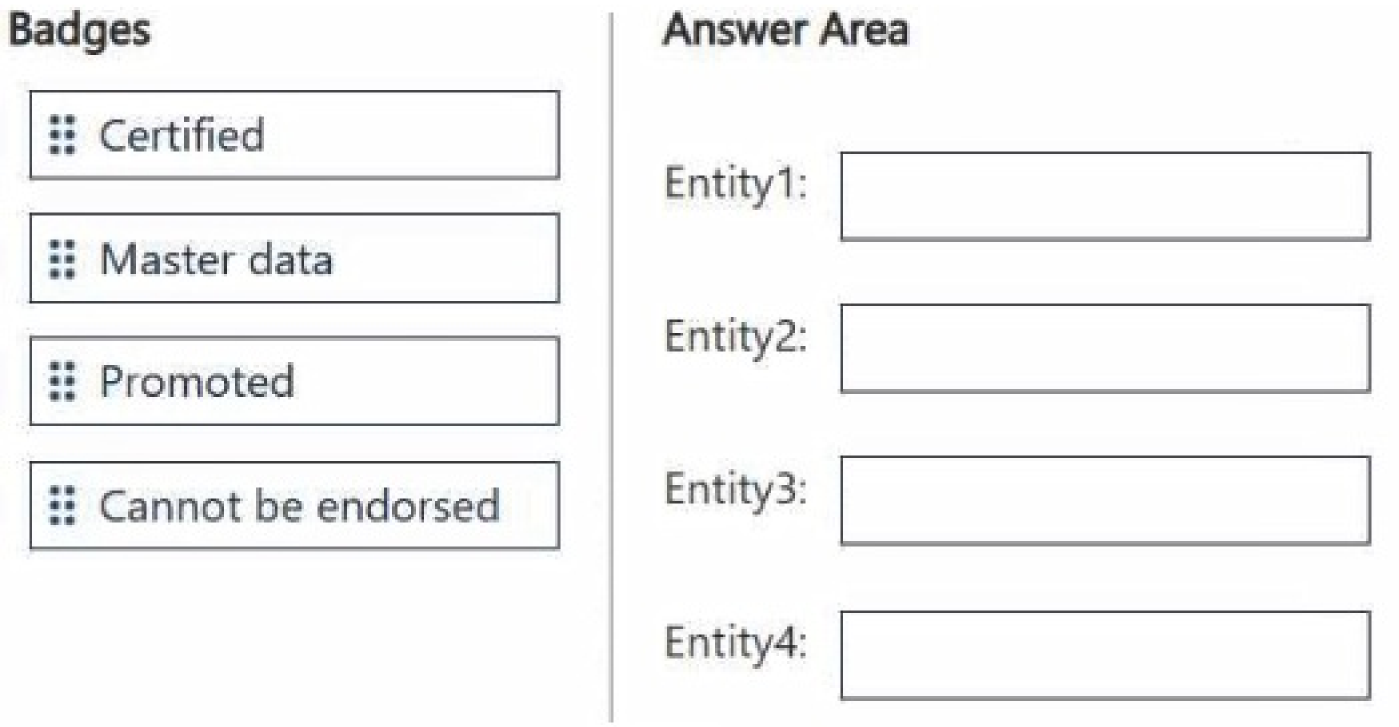

Answer:

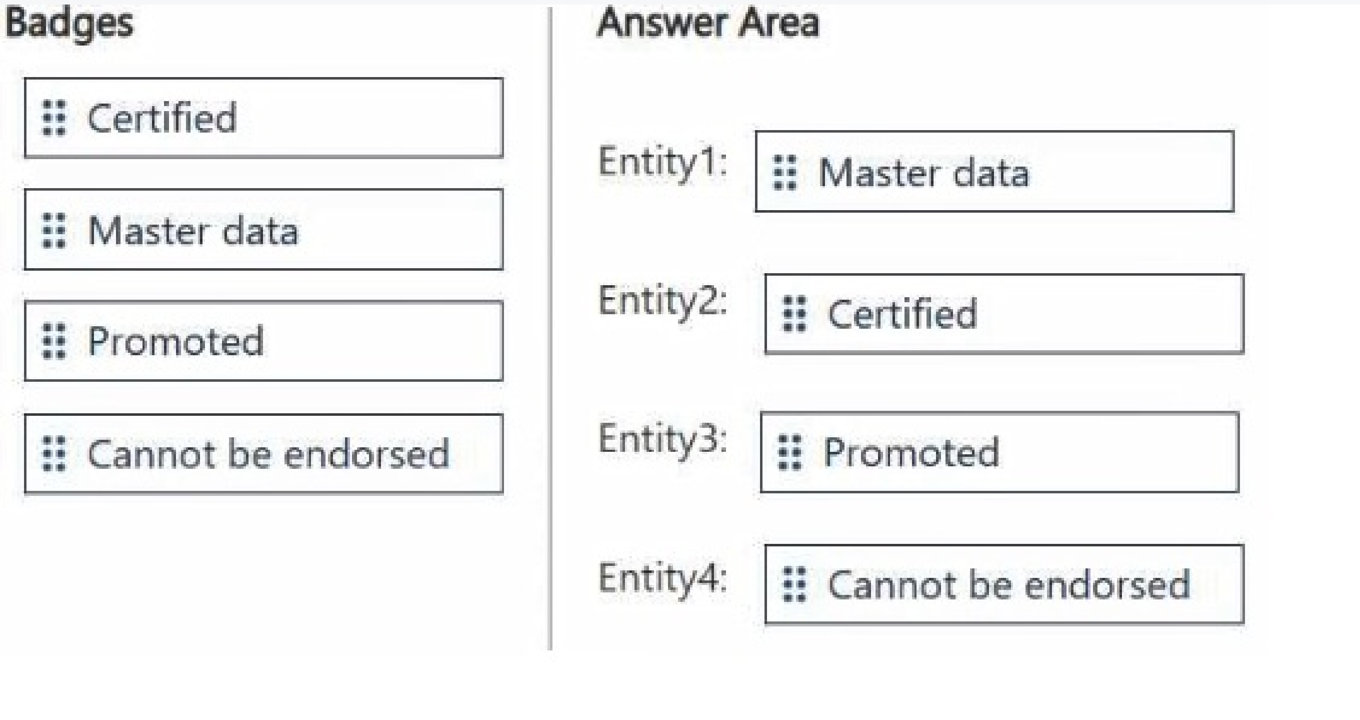

Explanation:

1.**Master Data.**

Refers to authoritative data that is central to business operations, often stored in a master data management system.

This is typically well-maintained and used across multiple departments.

Assigned to Entity1, as it represents centralized and validated business data.

2.**Certified.**

Indicates that an entity (such as a dataset or report) is officially validated by an authority in the organization.

Typically used for trusted and critical business data.

Assigned to Entity2 because this entity meets the highest quality standards.

3.**Promoted.**

Indicates that an entity is recommended for use but is not fully certified.

This badge is usually given when an item is considered useful but has not gone through a formal approval process.

Assigned to Entity3, which signifies that it is endorsed for use but not yet fully certified.

4. **Cannot be Endorsed.**

Indicates that an entity does not qualify for endorsement (either promoted or certified).

This could be due to low-quality data, lack of validation, or experimental datasets.

Assigned to Entity4, meaning it has not met the standards for endorsement.

Question: 24

Exam Author

HOTSPOT -
You have three users named User1, User2, and User3.
You have the Fabric workspaces shown in the following table.

Name	Workspace admin
Workspace1	User1
Workspace2	User2

You have a security group named Group1 that contains User1 and User3.
The Fabric admin creates the domains shown in the following table.

Name	Domain admin
Domain1	User1
Domain2	User2

User1 creates a new workspace named Workspace3.
You add Group1 to the default domain of Domain1.
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
User3 has Viewer role access to Workspace3.	○	○
User3 has Domain contributor access to Domain1.	○	○
User2 has Contributor role access to Workspace3.	○	○

Answer:

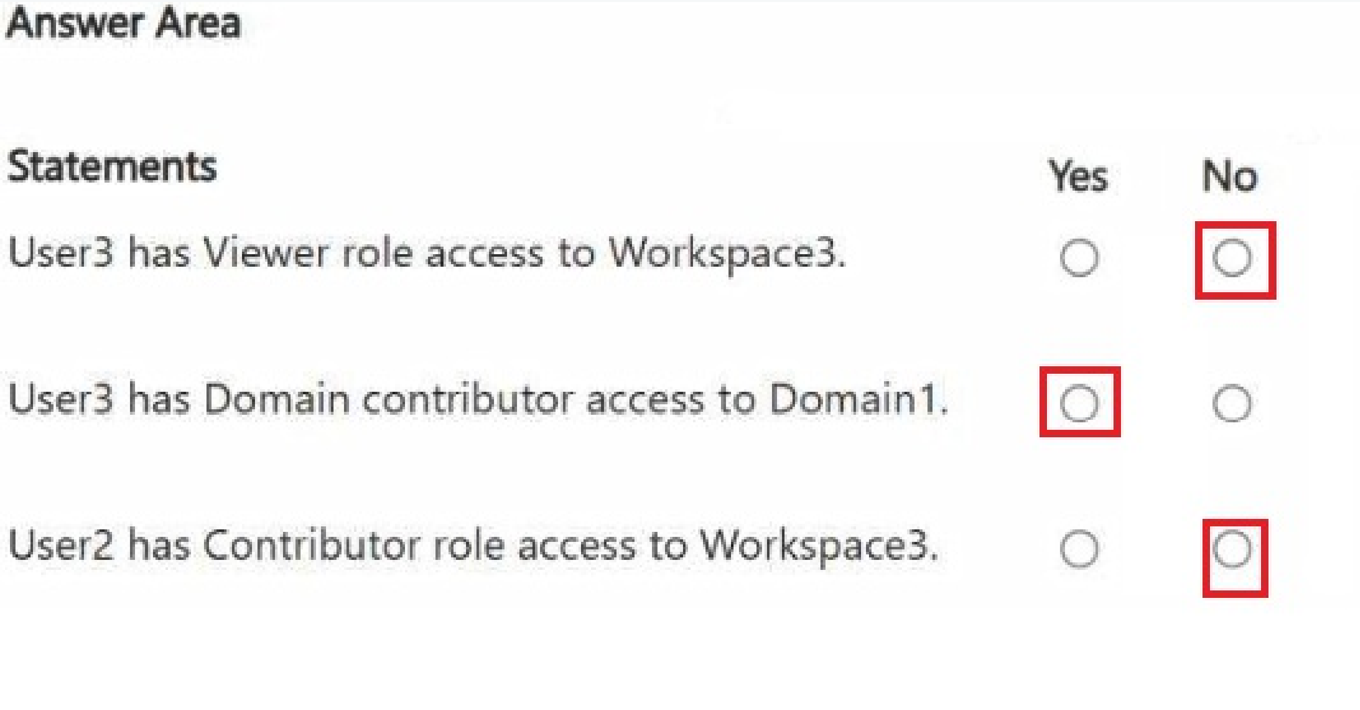

Explanation:

User3 has Viewer role access to Workspace3.

No

User3 has Domain Contributor access to Domain1.

The "**Yes**" option is selected, meaning User3 has Domain Contributor permissions in Domain1.

The Domain Contributor role typically allows managing content within a domain but does not grant full admin rights.

User2 has Contributor role access to Workspace3.

The **"No"** option is selected, meaning User2 does NOT have Contributor access to Workspace3.

The Contributor role would allow editing content in the workspace, but since "No" is selected, User2 lacks these permissions.

Question: 25 **Exam Author**

You have two Fabric workspaces named Workspace1 and Workspace2.
You have a Fabric deployment pipeline named deployPipeline1 that deploys items from Workspace1 to Workspace2. DeployPipeline1 contains all the items in Workspace1.
You recently modified the items in Workspaces1.
The workspaces currently contain the items shown in the following table.

Workspace	Items
Workspace1	Model1 Notebook1 Report1 Lakehouse1 Pipeline1
Workspace2	Model1 Notebook2 Report1 Lakehouse2

Items in Workspace1 that have the same name as items in Workspace2 are currently paired.
You need to ensure that the items in Workspace1 overwrite the corresponding items in Workspace2. The solution must minimize effort.
What should you do?

A.Delete all the items in Workspace2, and then run deployPipeline1.

B.Rename each item in Workspace2 to have the same name as the items in Workspace1.

C.Back up the items in Workspace2, and then run deployPipeline1.

D.Run deployPipeline1 without modifying the items in Workspace2.

Answer: D

Explanation:

D. Run deployPipeline1 without modifying the items in Workspace2.

When items in Workspace1 and Workspace2 are paired and you run the deployment pipeline (deployPipeline1), the pipeline will automatically update the paired items in Workspace2 with the changes made in Workspace1. This means that the modifications in Workspace1 will overwrite the corresponding items in Workspace2 without requiring any additional steps.

Question: 26

You have a Fabric workspace named Workspace1 that contains a data pipeline named Pipeline1 and a lakehouse named Lakehouse1.
You have a deployment pipeline named deployPipeline1 that deploys Workspace1 to Workspace2.
You restructure Workspace1 by adding a folder named Folder1 and moving Pipeline1 to Folder1.
You use deployPipeline1 to deploy Workspace1 to Workspace2.
What occurs to Workspace2?

A.Folder1 is created, Pipeline1 moves to Folder1, and Lakehouse1 is deployed.

B.Only Pipeline1 and Lakehouse1 are deployed.

C.Folder1 is created, and Pipeline1 and Lakehouse1 move to Folder1.

D.Only Folder1 is created and Pipeline1 moves to Folder1.

Answer: A

Explanation:

A. Folder1 is created, Pipeline1 moves to Folder1, and Lakehouse1 is deployed.

Folder1 is created: The deployment pipeline will replicate the structure of Workspace1 in Workspace2, including the creation of Folder1.

Pipeline1 moves to Folder1: Since Pipeline1 was moved to Folder1 in Workspace1, it will be deployed to Folder1 in Workspace2.

Lakehouse1 is deployed: Lakehouse1 is part of Workspace1 and will be deployed to Workspace2 as part of the deployment process.

Question: 27 Exam Author

DRAG DROP -
Your company has a team of developers. The team creates Python libraries of reusable code that is used to transform data.
You create a Fabric workspace name Workspace1 that will be used to develop extract, transform, and load (ETL) solutions by using notebooks.
You need to ensure that the libraries are available by default to new notebooks in Workspace1.
Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

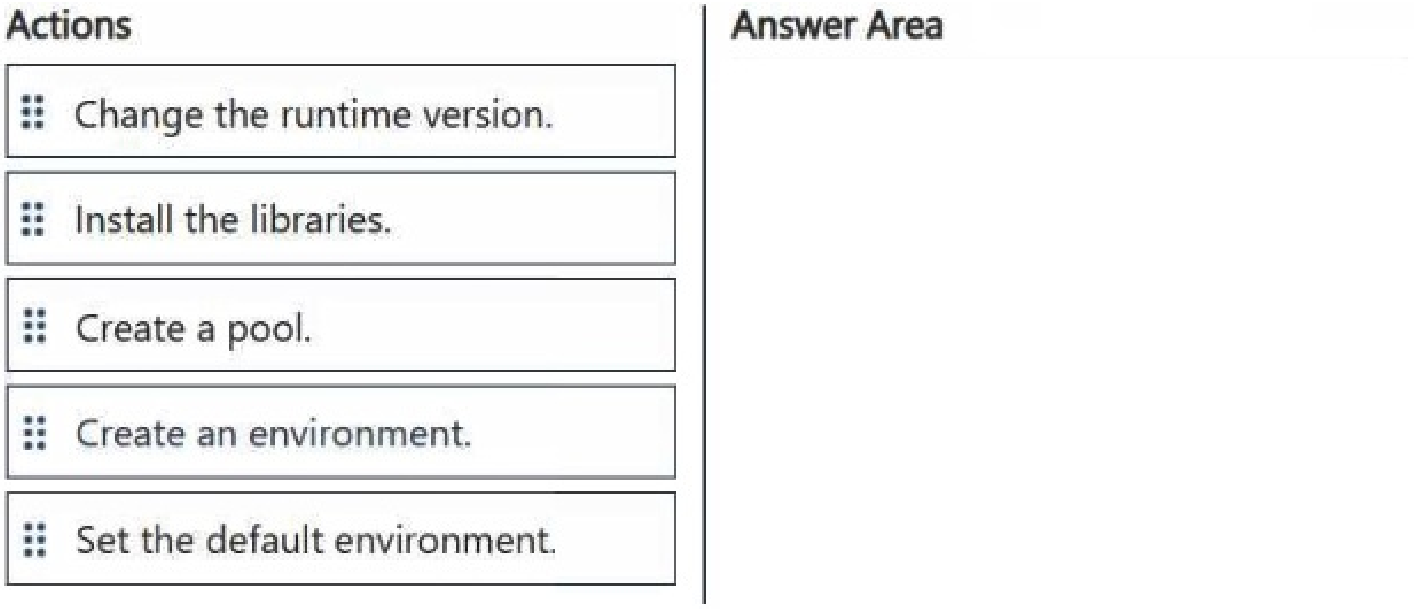

Answer:

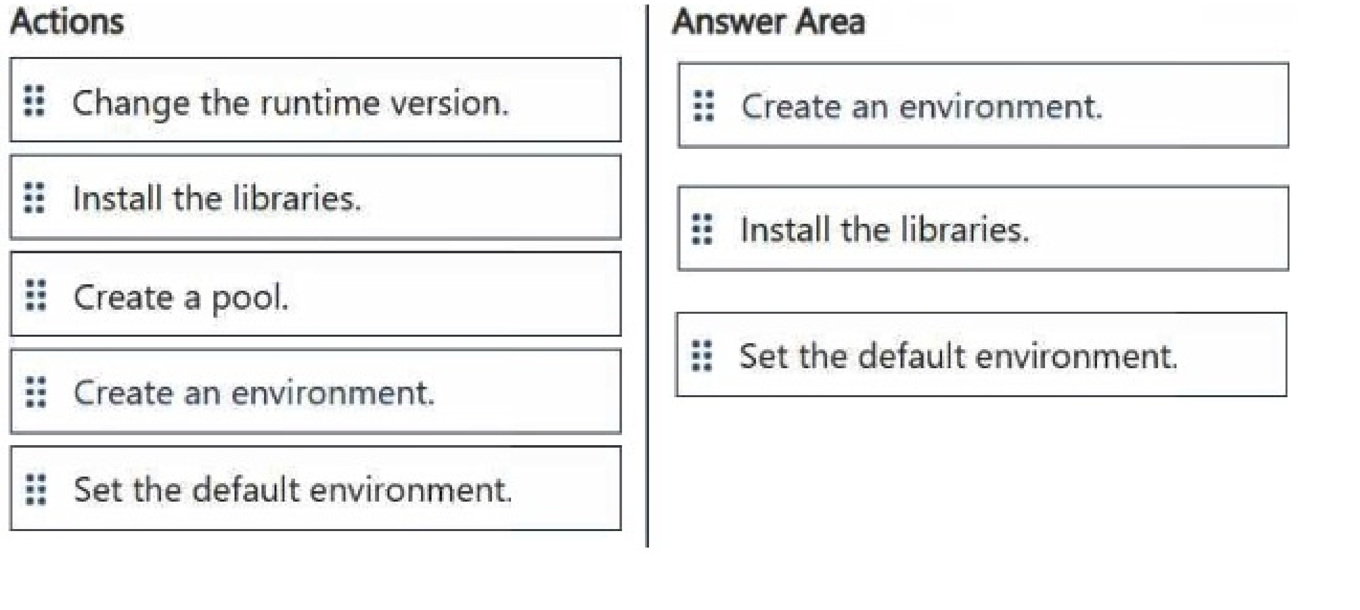

Explanation:

Create an environment: Matches with "Create an environment."

This action involves defining an environment where libraries, dependencies, and configurations are managed.

Install the libraries :Matches with "Install the libraries."

Installing libraries involves setting up necessary packages required for development and execution.

Set the default environment : Matches with "Set the default environment."

This action defines a specific environment as the default for execution.

Question: 28

You have a Fabric workspace that contains a lakehouse and a notebook named Notebook1. Notebook1 reads data into a DataFrame from a table named Table1 and applies transformation logic. The data from the DataFrame is then written to a new Delta table named Table2 by using a merge operation.
You need to consolidate the underlying Parquet files in Table1.
Which command should you run?

A.VACUUM
B.BROADCAST
C.OPTIMIZE
D.CACHE

Answer: C

Explanation:

OPTIMIZE: This command is used to compact small files into larger ones and optimize the layout of data in a Delta table. By running the OPTIMIZE command on Table1, you can consolidate the Parquet files and improve the performance of read and write operations on the table. To consolidate the underlying Parquet files in Table1, you should run the OPTIMIZE command.

Question: 29 **Exam Author**

You have five Fabric workspaces.
You are monitoring the execution of items by using Monitoring hub.
You need to identify in which workspace a specific item runs.
Which column should you view in Monitoring hub?

A.Start time

B.Capacity

C.Activity name

D.Submitter

E.Item type

F.Job type

G.Location

Answer: G

Explanation:

The Location shows the Workspace.

Location: This column displays the workspace where the item is being executed, helping you pinpoint the exact workspace of the item.

Reference:

https://learn.microsoft.com/en-us/training/modules/monitor-fabric-items/3-use-monitor-hub

Question: 30 **Exam Author**

You have a Fabric workspace that contains a warehouse named DW1. DW1 is loaded by using a notebook named Notebook1.
You need to identify which version of Delta was used when Notebook1 was executed.
What should you use?

A.Real-Time hub

B.OneLake data hub

C.the Admin monitoring workspace

D.Fabric Monitor

E.the Microsoft Fabric Capacity Metrics app

Answer: D

Explanation:

D. Fabric Monitor.

Fabric Monitor: This tool provides detailed monitoring and logging capabilities for various components within a Fabric workspace, including notebooks and data processing tasks. By using Fabric Monitor, you can track and analyze the execution details of Notebook1, including the version of Delta used during its execution. This information is crucial for debugging, auditing, and ensuring compatibility across different versions of Delta.

Question: 31

Exam Author

DRAG DROP -

You have a Fabric workspace that contains a warehouse named Warehouse1.

In Warehouse1, you create a table named DimCustomer by running the following statement.

```
CREATE TABLE dbo.DimCustomer (
  CustomerKey VARCHAR(255) NOT NULL,
  Name VARCHAR(255) NOT NULL,
  Email VARCHAR(255) NOT NULL
  );
```

You need to set the Customerkey column as a primary key of the DimCustomer table.

Which three code segments should you run in sequence? To answer, move the appropriate code segments from the list of code segments to the answer area and arrange them in the correct order.

Code Segments

- DROP CONSTRAINT PK_DimCustomer
- ADD CONSTRAINT PK_DimCustomer PRIMARY KEY NONCLUSTERED (CustomerKey)
- NOT ENFORCED
- ALTER TABLE dbo.DimCustomer
- ADD CONSTRAINT PK_DimCustomer PRIMARY KEY CLUSTERED (CustomerKey)
- ENFORCED

Answer Area

Answer:

Code Segments

- DROP CONSTRAINT PK_DimCustomer
- ADD CONSTRAINT PK_DimCustomer PRIMARY KEY NONCLUSTERED (CustomerKey)
- NOT ENFORCED
- ALTER TABLE dbo.DimCustomer
- ADD CONSTRAINT PK_DimCustomer PRIMARY KEY CLUSTERED (CustomerKey)
- ENFORCED

Answer Area

- ALTER TABLE dbo.DimCustomer
- ADD CONSTRAINT PK_DimCustomer PRIMARY KEY NONCLUSTERED (CustomerKey)
- NOT ENFORCED

Explanation:

ALTER TABLE dbo.DimCustomer.

This is necessary to modify the structure of an existing table.

Since adding or dropping a primary key constraint requires modifying a table, this statement is correct.

ADD CONSTRAINT PK_DimCustomer PRIMARY KEY NONCLUSTERED (CustomerKey)

This statement is used to define a primary key on the CustomerKey column.

It specifies a NONCLUSTERED primary key, meaning the physical ordering of data is not changed, and a separate index structure is created.

This selection aligns with the requirement of having a nonclustered primary key.

NOT ENFORCED

In some data warehousing scenarios, constraints might not be enforced to allow better query performance and faster data ingestion.

If the system does not enforce referential integrity (e.g., in Azure Synapse Analytics), this would be applicable.

Question: 32 **Exam Author**

You have a Fabric workspace that contains a semantic model named Model1.
You need to dynamically execute and monitor the refresh progress of Model1.
What should you use?

A.dynamic management views in Microsoft SQL Server Management Studio (SSMS)
B.Monitoring hub
C.dynamic management views in Azure Data Studio
D.a semantic link in a notebook

Answer: D

Explanation:

D. a semantic link in a notebook.

Semantic link in a notebook: This approach allows you to dynamically execute operations and monitor the refresh progress of the semantic model (Model1) within the interactive and flexible environment of a notebook. By using a semantic link, you can write custom scripts to trigger the refresh process and track its progress in real-time. This method provides a high degree of control and visibility over the operations on your semantic model.

Question: 33 **Exam Author**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have a Fabric eventstream that loads data into a table named Bike_Location in a KQL database. The table contains the following columns:
BikepointID -
Street -
Neighbourhood -
No_Bikes -
No_Empty_Docks -
Timestamp -
You need to apply transformation and filter logic to prepare the data for consumption. The solution must return data for a neighbourhood named Sands End when No_Bikes is at least 15. The results must be ordered by No_Bikes

in ascending order.
Solution: You use the following code segment:

```
bike_location
| filter Neighbourhood == "Sands End" and No_Bikes >= 15
| sort by No_Bikes
| project BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp
| project BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp
```

Does this meet the goal?

A.Yes
B.No

Answer: B

Explanation:

B. No

If you want, you can paste the exact code segment and I'll point out precisely which line breaks the requirement.

Question: 34

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have a Fabric eventstream that loads data into a table named Bike_Location in a KQL database. The table contains the following columns:
BikepointID -
Street -
Neighbourhood -
No_Bikes -
No_Empty_Docks -
Timestamp -
You need to apply transformation and filter logic to prepare the data for consumption. The solution must return data for a neighbourhood named Sands End when No_Bikes is at least 15. The results must be ordered by No_Bikes in ascending order.
Solution: You use the following code segment:

```
bike_location
| filter Neighbourhood == "Sands End" and No_Bikes >= 15
| order by No_Bikes
| project BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp
```

Does this meet the goal?

A.Yes
B.No

Answer: B

Explanation:

The default sorting order in KQL is **descending** (desc), not ascending (asc).

The solution does not explicitly specify asc in the order by clause, so the results will be sorted in descending order by default.

The requirement is to sort the data by No_Bikes in **ascending order**, which is not achieved without explicitly specifying asc.

Why other answers are not correct:

A. Yes: This would be incorrect because the solution fails to meet the requirement of sorting in ascending order due to the default descending behavior in KQL.

Important Tip:

Always explicitly specify the sorting order (asc or desc) in KQL to avoid confusion, especially since its default behavior differs from SQL.

Question: 35 Exam Author

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have a Fabric eventstream that loads data into a table named Bike_Location in a KQL database. The table contains the following columns:
BikepointID -
Street -
Neighbourhood -
No_Bikes -
No_Empty_Docks -
Timestamp -
You need to apply transformation and filter logic to prepare the data for consumption. The solution must return data for a neighbourhood named Sands End when No_Bikes is at least 15. The results must be ordered by No_Bikes in ascending order.
Solution: You use the following code segment:

```
bike_location
| filter Neighbourhood == "Sands End" and No_Bikes >= 15
| sort by No_Bikes asc
| project BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp
```

Does this meet the goal?

A.Yes
B.No

Answer: A

Explanation:

The sort and order operators are equivalent.

The provided code segment correctly filters the data for the neighborhood "Sands End" where the number of bikes (No_Bikes) is at least 15. It then explicitly sorts the results by No_Bikes in ascending order using sort by No_Bikes asc and projects the required columns (BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp). This meets all the stated goals of the problem.

Why other answers are not correct:

B. No: This would be incorrect because the solution explicitly specifies asc in the sort by clause, ensuring the data is ordered by No_Bikes in ascending order as required.

Important Tip:

Always ensure that the sorting order is explicitly specified in KQL to match the requirements, as the default behavior might differ from other query languages like SQL.

Reference:

https://learn.microsoft.com/en-us/kusto/query/sort-operator?view=microsoft-fabric

Question: 36 **Exam Author**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have a Fabric eventstream that loads data into a table named Bike_Location in a KQL database. The table contains the following columns:
BikepointID -
Street -
Neighbourhood -
No_Bikes -
No_Empty_Docks -
Timestamp -
You need to apply transformation and filter logic to prepare the data for consumption. The solution must return data for a neighbourhood named Sands End when No_Bikes is at least 15. The results must be ordered by No_Bikes in ascending order.
Solution: You use the following code segment:

```
SELECT BikepointID, Street, Neighbourhood, No_Bikes, No_Empty_Docks, Timestamp
FROM bike_location
WHERE neighbourhood = 'Sands End'
AND no_bikes >= 15
ORDER BY no_bikes
```

Does this meet the goal?

A.Yes
B.No

Answer: A

Explanation:

You can use SQL code directly without any additional "--" now (checked in Fabric). Also, it is case insensitive to the column and table names, order by sorts ascending by default. SQL By default shows ASC order.

Case Study -
This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.
To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.
At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.
To start the case study -
To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.
Overview -
Litware, Inc. is a publishing company that has an online bookstore and several retail bookstores worldwide. Litware also manages an online advertising business for the authors it represents.
Existing Environment. Fabric Environment
Litware has a Fabric workspace named Workspace1. High concurrency is enabled for Workspace1.
The company has a data engineering team that uses Python for data processing.
Existing Environment. Data Processing
The retail bookstores send sales data at the end of each business day, while the online bookstore constantly provides logs and sales data to a central enterprise resource planning (ERP) system.
Litware implements a medallion architecture by using the following three layers: bronze, silver, and gold. The sales data is ingested from the ERP system as Parquet files that land in the Files folder in a lakehouse. Notebooks are used to transform the files in a Delta table for the bronze and silver layers. The gold layer is in a warehouse that has V-Order disabled.
Litware has image files of book covers in Azure Blob Storage. The files are loaded into the Files folder.
Existing Environment. Sales Data
Month-end sales data is processed on the first calendar day of each month. Data that is older than one month never changes.
In the source system, the sales data refreshes every six hours starting at midnight each day.
The sales data is captured in a Dataflow Gen1 dataflow. When the dataflow runs, new and historical data is captured. The dataflow captures the following fields of the source:
•Sales Date
•Author
•Price
•Units
•SKU
A table named AuthorSales stores the sales data that relates to each author. The table contains a column named AuthorEmail. Authors authenticate to a guest Fabric tenant by using their email address.
Existing Environment. Security Groups
Litware has the following security groups:
•Sales
•Fabric Admins
•Streaming Admins
Existing Environment. Performance Issues
Business users perform ad-hoc queries against the warehouse. The business users indicate that reports against the warehouse sometimes run for two hours and fail to load as expected. Upon further investigation, the data engineering team receives the following error message when the reports fail to load: “The SQL query failed while running.”
The data engineering team wants to debug the issue and find queries that cause more than one failure.
When the authors have new book releases, there is often an increase in sales activity. This increase slows the data ingestion process.
The company’s sales team reports that during the last month, the sales data has NOT been up-to-date when they arrive at work in the morning.
Requirements. Planned Changes -
Litware recently signed a contract to receive book reviews. The provider of the reviews exposes the data in Amazon Simple Storage Service (Amazon S3) buckets.
Litware plans to manage Search Engine Optimization (SEO) for the authors. The SEO data will be streamed from a REST API.
Requirements. Version Control -

Litware plans to implement a version control solution in Fabric that will use GitHub integration and follow the principle of least privilege.
Requirements. Governance Requirements
To control data platform costs, the data platform must use only Fabric services and items. Additional Azure resources must NOT be provisioned.
Requirements. Data Requirements -
Litware identifies the following data requirements:
•Process the SEO data in near-real-time (NRT).
•Make the book reviews available in the lakehouse without making a copy of the data.
•When a new book cover image arrives in the Files folder, process the image as soon as possible.
You need to ensure that processes for the bronze and silver layers run in isolation.
How should you configure the Apache Spark settings?

A.Disable high concurrency.

B.Create a custom pool.

C.Modify the number of executors.

D.Set the default environment.

Answer: B

Explanation:

B. Create a custom pool.

While disabling high concurrency (Option A) might seem like it isolates processes, it's not the recommended approach for managing isolation in layered architectures like bronze and silver. By creating a custom pool (Option B), you can allocate dedicated resources to each layer, ensuring they run independently without interfering with one another. Custom pools give you fine-grained control over resource allocation, making them the ideal solution for this scenario.

Question: 38

DRAG DROP
-
Case Study
-
This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.
To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.
At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.
To start the case study
-
To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.
Overview
-
Litware, Inc. is a publishing company that has an online bookstore and several retail bookstores worldwide. Litware also manages an online advertising business for the authors it represents.
Existing Environment. Fabric Environment
Litware has a Fabric workspace named Workspace1. High concurrency is enabled for Workspace1.

The company has a data engineering team that uses Python for data processing.
Existing Environment. Data Processing
The retail bookstores send sales data at the end of each business day, while the online bookstore constantly provides logs and sales data to a central enterprise resource planning (ERP) system.
Litware implements a medallion architecture by using the following three layers: bronze, silver, and gold. The sales data is ingested from the ERP system as Parquet files that land in the Files folder in a lakehouse. Notebooks are used to transform the files in a Delta table for the bronze and silver layers. The gold layer is in a warehouse that has V-Order disabled.
Litware has image files of book covers in Azure Blob Storage. The files are loaded into the Files folder.
Existing Environment. Sales Data
Month-end sales data is processed on the first calendar day of each month. Data that is older than one month never changes.
In the source system, the sales data refreshes every six hours starting at midnight each day.
The sales data is captured in a Dataflow Gen1 dataflow. When the dataflow runs, new and historical data is captured. The dataflow captures the following fields of the source:
•Sales Date
•Author
•Price
•Units
•SKU
A table named AuthorSales stores the sales data that relates to each author. The table contains a column named AuthorEmail. Authors authenticate to a guest Fabric tenant by using their email address.
Existing Environment. Security Groups
Litware has the following security groups:
•Sales
•Fabric Admins
•Streaming Admins
Existing Environment. Performance Issues
Business users perform ad-hoc queries against the warehouse. The business users indicate that reports against the warehouse sometimes run for two hours and fail to load as expected. Upon further investigation, the data engineering team receives the following error message when the reports fail to load: “The SQL query failed while running.”
The data engineering team wants to debug the issue and find queries that cause more than one failure.
When the authors have new book releases, there is often an increase in sales activity. This increase slows the data ingestion process.
The company’s sales team reports that during the last month, the sales data has NOT been up-to-date when they arrive at work in the morning.
Requirements. Planned Changes
-
Litware recently signed a contract to receive book reviews. The provider of the reviews exposes the data in Amazon Simple Storage Service (Amazon S3) buckets.
Litware plans to manage Search Engine Optimization (SEO) for the authors. The SEO data will be streamed from a REST API.
Requirements. Version Control
-
Litware plans to implement a version control solution in Fabric that will use GitHub integration and follow the principle of least privilege.
Requirements. Governance Requirements
To control data platform costs, the data platform must use only Fabric services and items. Additional Azure resources must NOT be provisioned.
Requirements. Data Requirements
Litware identifies the following data requirements:
•Process the SEO data in near-real-time (NRT).
•Make the book reviews available in the lakehouse without making a copy of the data.
•When a new book cover image arrives in the Files folder, process the image as soon as possible.
You need to ensure that the authors can see only their respective sales data.
How should you complete the statement? To answer, drag the appropriate values the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Values

- AuthorSales
- AuthorEmail
- AuthorSales.AuthorEmail
- BLOCK
- FILTER
- INLINE
- SCHEMABINDING
- USER_NAME()

Answer Area

```
CREATE FUNCTION dbo.tvf_rlspredicate(@Author AS varchar(50))
    RETURNS TABLE
WITH [          ]
AS
    RETURN SELECT 1 AS tvf_rlspredicate_result
WHERE @Author = [          ]
GO

CREATE SECURITY POLICY RLSFilter
ADD FILTER PREDICATE Security.tvf_rlspredicate(AuthorEmail)
ON [          ]
WITH (STATE = ON)
```

Answer:

Answer Area

```
CREATE FUNCTION dbo.tvf_rlspredicate(@Author AS varchar(50))
    RETURNS TABLE
WITH SCHEMABINDING
AS
    RETURN SELECT 1 AS tvf_rlspredicate_result
WHERE @Author = USER_NAME()
GO

CREATE SECURITY POLICY RLSFilter
ADD FILTER PREDICATE Security.tvf_rlspredicate(AuthorEmail)
ON AuthorSales
WITH (STATE = ON)
```

Explanation:

SCHEMABINDING: Ensures the function is bound to the schema of the referenced objects. Required for RLS functions.

USER_NAME() returns the SQL Server login or database user name.

AuthorSales: The RLS policy is applied to the AuthorSales table.

Question: 39 **Exam Author**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have an Azure key vault named KeyVault1 that contains secrets.
You have a Fabric workspace named Workspace1. Workspace contains a notebook named Notebook1 that performs the following tasks:
•Loads stage data to the target tables in a lakehouse
•Triggers the refresh of a semantic model
You plan to add functionality to Notebook1 that will use the Fabric API to monitor the semantic model refreshes.
You need to retrieve the registered application ID and secret from KeyVault1 to generate the authentication token.
Solution: You use the following code segment:
Use notebookutils.credentials.getSecret and specify the key vault URL and key vault secret.
Does this meet the goal?

A.Yes
B.No

Answer: B

Explanation:

The method notebookutils.credentials.getSecret() in Microsoft Fabric does not accept a Key Vault URL. Instead, it requires the name of a linked service (which securely points to the Key Vault) and the name of the secret.

Question: 40 **Exam Author**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have an Azure key vault named KeyVault1 that contains secrets.
You have a Fabric workspace named Workspace1. Workspace contains a notebook named Notebook1 that performs the following tasks:
•Loads stage data to the target tables in a lakehouse
•Triggers the refresh of a semantic model
You plan to add functionality to Notebook1 that will use the Fabric API to monitor the semantic model refreshes.
You need to retrieve the registered application ID and secret from KeyVault1 to generate the authentication token.
Solution: You use the following code segment:
Use notebookutils.credentials.putSecret and specify the key vault URL and key vault secret.
Does this meet the goal?

A.Yes
B.No

Answer: B

Explanation:

You need to retrieve the registered application ID and secret from KeyVault1 to generate the authentication token. The function notebookutils.credentials.putSecret is used to store a secret into a secret scope — not retrieve it.

Question: 41 **Exam Author**

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have an Azure key vault named KeyVault1 that contains secrets.
You have a Fabric workspace named Workspace1. Workspace contains a notebook named Notebook1 that performs the following tasks:
•Loads stage data to the target tables in a lakehouse
•Triggers the refresh of a semantic model
You plan to add functionality to Notebook1 that will use the Fabric API to monitor the semantic model refreshes.
You need to retrieve the registered application ID and secret from KeyVault1 to generate the authentication token.
Solution: You use the following code segment:
Use notebookutils.credentials.getSecret and specify the key vault URL and the name of a linked service.
Does this meet the goal?

A.Yes
B.No

Answer: B

Explanation:

No, this solution does not meet the goal. The notebookutils.credentials.getSecret method requires the key vault URL and the name of the secret, not the name of a linked service.

Question: 42 **Exam Author**

DRAG DROP
-
You have two Fabric notebooks named Load_Salesperson and Load_Orders that read data from Parquet files in a lakehouse. Load_Salesperson writes to a Delta table named dim_salesperson. Load_Orders writes to a Delta table named fact_orders and is dependent on the successful execution of Load_Salesperson.
You need to implement a pattern to dynamically execute Load_Salesperson and Load_Orders in the appropriate order by using a notebook.
How should you complete the code? To answer, drag the appropriate values the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Values

- activities
- broadcast
- dependencies
- execute
- notebooks
- runMultiple

Answer Area

```
DAG = {
    "            ": [
      {
          "name": "Load_Salesperson",
          "path": "Load_Salesperson",
          "timeoutPerCellInSeconds": 300,
      },
      {
          "name": "Load_Orders",
          "path": "Load_Orders",
          "timeoutPerCellInSeconds": 600,
          "            ": ["Load_Salesperson"]
      }
    ],
    "timeoutInSeconds": 43200
}
mssparkutils.notebook.            (DAG)
```

Answer:

Answer Area

```
DAG = {
    "activities": [
        {
            "name": "Load_Salesperson",
            "path": "Load_Salesperson",
            "timeoutPerCellInSeconds": 300,
        },
        {
            "name": "Load_Orders",
            "path": "Load_Orders",
            "timeoutPerCellInSeconds": 600,
            "dependencies": ["Load_Salesperson"]
        }
    ],
    "timeoutInSeconds": 43200
}
mssparkutils.notebook.runMultiple(DAG)
```

Explanation:

activities: to list the activities.

dependencies: to make sure the order and dependency of running.

run Multiple: there are 2 notebooks to be executed.

Question: 43

Exam Author

HOTSPOT

-

You have a Fabric workspace named Workspace1 that contains a warehouse named Warehouse2.
A team of data analysts has Viewer role access to Workspace1.
You create a table by running the following statement.

```
CREATE TABLE [warehouse2].[dbo].[CreditCard]
(
  CreditCard varchar(20) NOT NULL
  ,CreditCardType varchar(10) NOT NULL)
GO
```

You need to ensure that the team can view only the first two characters and the last four characters of the CreditCard attribute.
How should you complete the statement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

[dropdown] `TABLE dbo.CreditCard`

- ALTER
- CREATE
- DEFAULT
- DROP
- EMAIL
- PARTIAL
- REPLACE
- UPDATE

[dropdown] `COLUMN [CreditCard]`

- ALTER
- CREATE
- DEFAULT
- DROP
- EMAIL
- PARTIAL
- REPLACE
- UPDATE

`ADD MASKED`

`WITH (FUNCTION = '` [dropdown] `(2,"XXXXXXXXXX",4)')`

- ALTER
- CREATE
- DEFAULT
- DROP
- EMAIL
- PARTIAL
- REPLACE
- UPDATE

Answer:

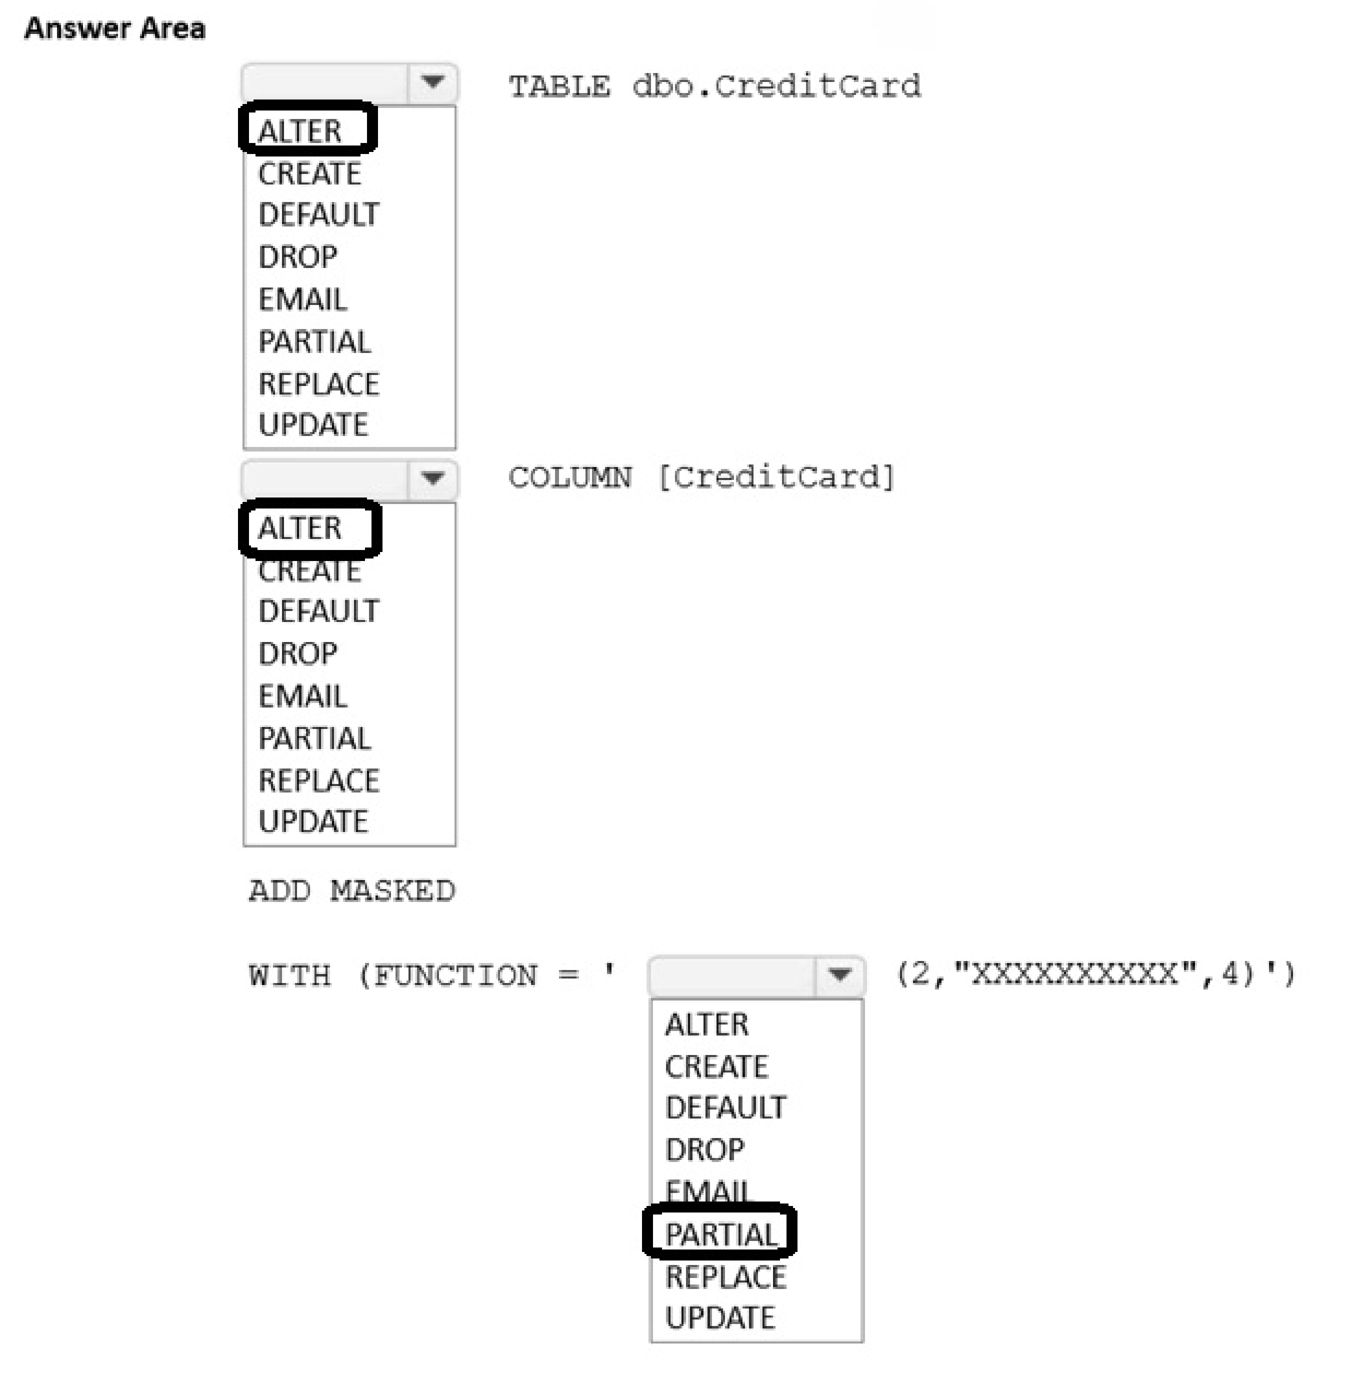

Explanation:

ALTER TABLE- to modify the tabel

ALTER COLUMN - to modify the column

PARTIAL(prefix padding, padding string, suffix padding) - to expose first and last n characters, adding custom padding 'xxx' of a text field